

IN4343 – Real Time Systems

June 23rd 2016, from 13:30 to 16:30

Koen Langendoen

Question:	1	2	3	4	5	6	Total
Points:	15	15	15	25	10	10	90
Score:							

- This is a closed book exam
- You may use a **simple** calculator only (i.e. graphical calculators are not permitted)
- Write your answers with a black or blue pen, not with a pencil
- **Always justify your answers, unless stated otherwise**
- The exam covers the following material:
 - (a) chapters 1-6, 8-9 of the book “Hard Real-Time Computing Systems (3rd ed)” by G. Buttazzo
 - (b) the paper “The Worst-Case Execution-Time Problem” by Wilhelm et al. (except Section 6)
 - (c) the paper “Transforming Execution-Time Boundable Code into Temporally Predictable Code” by P. Puschner
 - (d) the paper “Best-case response times and jitter analysis of real-time tasks” by R.J. Bril, E.F.M. Steffens, and W.F.J. Verhaegh

Liu and Layland (LL) bound		$U_{lub}^{RM} = n(2^{1/n} - 1)$
Hyperbolic (HB) bound		$\prod_{i=1}^n (U_i + 1) \leq 2$
Response Time Analysis		$WR_i = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{WR_i + AJ_k}{T_k} \right\rceil C_k$ $BR_i = C_i + \sum_{k=1}^{i-1} \left(\left\lceil \frac{BR_i - AJ_k}{T_k} \right\rceil - 1 \right)^+ C_k$ $w^+ = \max(w, 0)$
Processor Demand	schedulability	$g(t_1, t_2) = \sum_{r_i \geq t_1}^{d_i \leq t_2} C_i \quad g(0, L) = \sum_{i=1}^n \left\lfloor \frac{L + T_i - D_i}{T_i} \right\rfloor C_i$ $\forall L \in D, \quad g(0, L) \leq L$ $D = \{d_k d_k \leq \min(H, \max(D_{max}, L^*))\}$ $H = lcm(T_1, \dots, T_n)$ $L^* = \frac{\sum_{i=1}^n (T_i - D_i) U_i}{1 - U}$
Polling Server	schedulability	$U_{lub}^{RM+PS} = U_s + n \left[\left(\frac{2}{U_s + 1} \right)^{1/n} - 1 \right]$ $\prod_{i=1}^n (U_i + 1) \leq \frac{2}{U_s + 1}$
	response time	$R_a = C_a + \Delta_a + F_a(T_s - C_s)$ $\Delta_a = \left\lceil \frac{r_a}{T_s} \right\rceil T_s - r_a \quad F_a = \left\lceil \frac{C_a}{C_s} \right\rceil - 1$
Deferrable Server	schedulability	$U_{lub}^{RM+DS} = U_s + n \left[\left(\frac{U_s + 2}{2U_s + 1} \right)^{1/n} - 1 \right]$ $\prod_{i=1}^n (U_i + 1) \leq \frac{U_s + 2}{2U_s + 1}$
	response time	$R_a = C_a + \Delta_a - C_0 + F_a(T_s - C_s)$ $C_0 = \min(C_s(r_a), \Delta_a)$ $\Delta_a = \left\lceil \frac{r_a}{T_s} \right\rceil T_s - r_a \quad F_a = \left\lceil \frac{C_a - C_0}{C_s} \right\rceil - 1$
NP scheduling	level-i busy period	$L_i = B_i + \sum_{h=1}^i \left\lceil \frac{L_i}{T_h} \right\rceil C_h \quad N_i = \left\lceil \frac{L_i}{T_i} \right\rceil$ $B_i = \max_{j > i} \{C_j\}$
	response time	$s_{ik} = B_i + (k - 1)C_i + \sum_{h=1}^{i-1} \left(\left\lceil \frac{s_{ik}}{T_h} \right\rceil + 1 \right) C_h$ $R_{ik} = (s_{ik} + C_i) - (k - 1)T_i$ $R_i = \max_{k \in [1, N_i]} \{R_{ik}\}$
Elastic Model	utilization	$\forall i \quad U_i = U_{i0} - (U_0 - U_d) \frac{E_i}{E_S} \quad \text{where } E_S = \sum_{i=1}^n E_i$

Question 1

[15 points]

Bratley's algorithm is part of the class of scheduling algorithms handling aperiodic tasks. It is a heuristic algorithm that prunes (large) parts of the search tree.

- (a) 3 points Bratley's algorithm can be classified as $(1|no-preem|L_{max})$ according to Graham's notation. Explain what the classification denotes.

Solution: Uniprocessor algorithm for non-preemptive tasks that minimizes the maximum lateness.

- (b) 4 points Explain how Bratley's algorithm works.

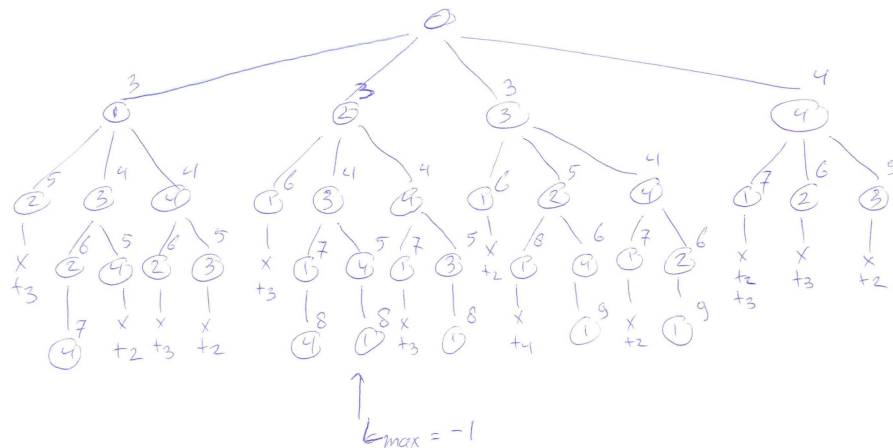
Solution: It expands the tree in depth-first order. A node, however, is only expanded when the partial schedule is found to be **strongly feasible**, i.e. if it remains feasible when adding any of the remaining nodes.

- (c) 8 points Apply Bratley's algorithm to find the feasible schedules for the following task set:

	a_i	C_i	d_i
τ_1	0	3	9
τ_2	1	2	6
τ_3	2	1	5
τ_4	3	1	8

Draw the resulting search tree, and determine the optimal L_{max} . Finally, compute the gain over expanding the tree fully.

Solution:



6 feasible schedules, with $\tau_2 - \tau_3 - \tau_4 - \tau_1$ achieving an $L_{max} = -1$. A fully expanded tree consists of $1 + 4 + 12 + 24 + 24 = 65$ nodes. Bratley only expands $1 + 4 + 12 + 12 + 6 = 35$ nodes, so the gain is $65/35 = 1.86$.

Question 2

[15 points]

	C_i	T_i
τ_1	2	5
τ_2	3	9
τ_3	3	15

(i)

	C_i	T_i
τ_1	2	5
τ_2	3	9
τ_3	4	15

(ii)

	C_i	T_i
τ_1	2	5
τ_2	3	9
τ_3	5	15

(iii)

- (a) 3 points Which task sets are feasible under Earliest Deadline First scheduling?


Solution:

- (i) Feasible as processor utilization $(42/45) \leq 1$
- (ii) Feasible as processor utilization $(45/45) \leq 1$
- (iii) Unfeasible as processor utilization $(48/45) > 1$

- (b) 7 points Which task sets are feasible under Rate Monotonic scheduling?

Solution:

(i) Feasible




- (ii) Unfeasible as (i) has no slack in the critical instant of task τ_3
- (iii) Unfeasible as EDF can't schedule it

- (c) 5 points Determine the maximum lateness for the given task sets under RM scheduling. Hint: standard techniques may not apply to unfeasible task sets.

Solution:

- (i) $L_{max} = 0$, see figure above
- (ii) Task τ_3 misses its first deadline, and continues running in the 2^{nd} period pushing that instance to the right. Thus simply looking at the critical instant is not enough. A safe bet is to check the completer hyper period, i.e. until $t=45$.



$$L_{max} = 4$$

(iii) As the system is overloaded L_{max} grows unbounded with time.

Question 3

[15 points]

	C_i	D_i	T_i
τ_1	2	3	5
τ_2	4	6	9
τ_3	4	30	40

Deadlines complicate matters. For EDF for example, we may need to revert to the processor demand criterion to determine the feasibility of a schedule. The processor demand is defined as

$$g(0, L) = \sum_{i=1}^n \left\lfloor \frac{L + T_i - D_i}{T_i} \right\rfloor C_i$$


- (a) 5 points Explain what $g(0, L)$ denotes, and why the term $T_i - D_i$ is part of the numerator.

Solution: $g(0, L)$ captures the amount of work that has to be executed by time L when starting at time 0. When counting the number of instances of task τ_i , we know that the demand goes up as soon as the next deadline is reached. This deadline occurs before the next period T_i , so we must bump the counter early by $T_i - D_i$. See the drawing in the lecture notes.

- (b) 10 points Determine if the task set is feasible under EDF scheduling.

Solution: Processor utilization is $340/360 \leq 1$, so we need to check the processor demand until $\min(H, \max(D_{max}, L^*))$ with $H = 360$, $D_{max} = 30$, and $L^* = 56$, so until 56.

	demand
$g(0,3)$	2
$g(0,6)$	6
$g(0,8)$	8
$g(0,13)$	10
$g(0,15)$	14
$g(0,18)$	16
$g(0,23)$	18
$g(0,24)$	22
$g(0,28)$	24
$g(0,30)$	28
$g(0,33)$	34



Task set is **unfeasible**!

Question 4

[25 points]

When mixing periodic and aperiodic tasks one can make use of a fixed priority server to schedule the aperiodic tasks. Consider the following periodic tasks and aperiodic jobs (under RM):

	C_i	T_i
τ_1	2	5
τ_2	2	10

	a_i	C_i
J_1	3	3
J_2	7	4

- (a) 4 points Dimension both a polling and a deferrable server by means of the hyperbolic bound guaranteeing schedulability.

Solution: $P = \prod_{i=1}^n (U_i + 1) = \frac{7}{5} \cdot \frac{12}{10} = \frac{84}{50}$

Polling server: $U_s = \frac{2 - P}{P} = \frac{4}{21}, < C_s = \frac{20}{21}, T_s = 5 >$

Deferrable server: $U_s = \frac{2 - P}{2P - 1} = \frac{8}{59}, < C_s = \frac{40}{59}, T_s = 5 >$

- (b) 7 points Consider a server dimensioned with $<C_s = 2, T_s = 5>$. Determine if this configuration is feasible (i.e. no deadline violations) to be used for a polling server. Determine the same for a deferrable server.

Solution: Polling server: utilization is 1 (!), we need to check the critical instant. Task τ_1 has a lateness of -1, task τ_2 has a lateness of 0, hence, it is a feasible task set.

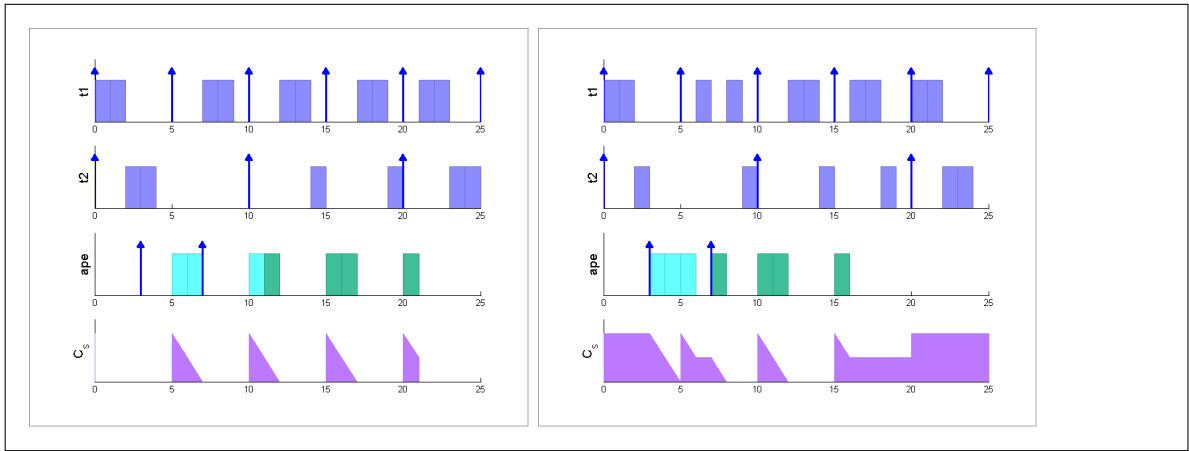
Deferrable server: utilization is 1 (!), we need to check the critical instant, but with jitter for the server ($AJ_s = 5 - 2 = 3$). As the schedule was already tight for the polling server, now tasks are getting into troubles ($R_1 = 6, R_2 = 20$) and violates their deadline. Thus the schedule is **unfeasible**.

- (c) 10 points Obtain the response times of jobs J_1 and J_2 for a polling server, as well as for a deferrable server when configured as $<C_s = 2, T_s = 5>$.

Solution: Response times

	PS	DS
J_1	8	3
J_2	14 (15*)	9

*the equations in the book for PS ignore the case that spare capacity can be used when a job is blocked by a predecessor, see picture below, so the computed response time is too high.



- (d) 4 points Determine the response times of the jobs when **slack stealing** would be used.

Solution: From the DS plot we can see that job J_1 is already optimal. Regarding job J_2 we can see that there is no slack until $t = 10$, but there is room in the 3rd period of task τ_1 and the 2nd period of τ_2 . Shifting both 1 unit to the right allows job J_2 to finish by $t = 13$, reducing its response time to 6.

Question 5

[10 points]

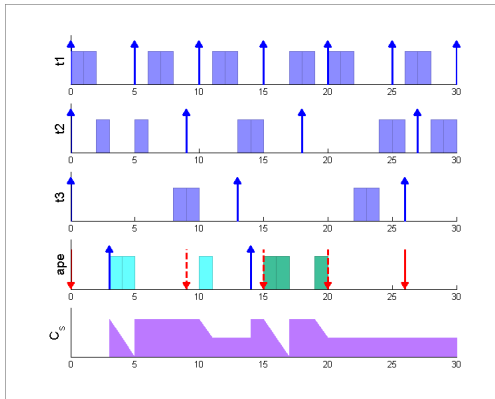
To protect periodic tasks from overruns by a-periodic jobs a Constant Bandwidth Server can be used in combination with an EDF scheduler.

	C_i	T_i
τ_1	2	5
τ_2	2	9
τ_3	2	13

	a_i	C_i
J_1	3	3
J_2	14	3

- (a) 7 points Consider the case that the CBS server is dimensioned as $\langle C_S = 2, T_S = 6 \rangle$, determine the response times of jobs J_1 and J_2 .

Solution:

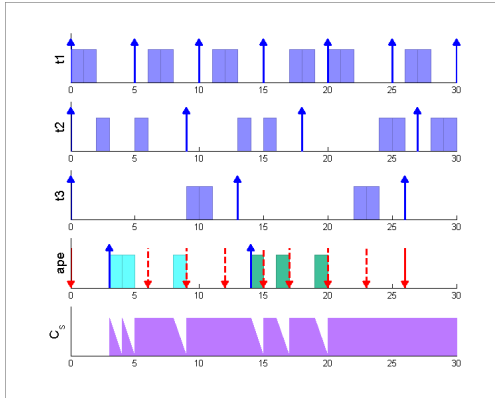


$$RJ_1 = 8, RJ_2 = 6$$

- (b) 3 points Argue the effect (if any) of keeping the server utilization constant, but halving its period (i.e. $C_S = 1, T_S = 3$).

Solution:

Things can only get better when the server period is reduced because there will be more intermediate deadlines (with in-between values) allowing shorter jobs to attain a higher priority (earlier deadline). Proof by example (not required for full score):



$$RJ_1 = 6, RJ_2 = 6$$

Question 6**[10 points]**

One approach to handle permanent overloads is to resort to value-based scheduling, in which the least important tasks are simply rejected. The performance of a scheduling algorithm A can be evaluated through its cumulative value $\Gamma_A(T) < \Gamma_{max}(T) = \sum_i V_i$ for task set T with V_i denoting the value of task i .

- (a) 4 points Show that, in case of overload, an **online** scheduling algorithm can not be optimal. That is, achieve the same performance as an offline (clairvoyant) algorithm in all cases.

Solution: Construct a counter example showing that knowledge of the future is required to decide to skip a task or not. See lecture notes for a concrete example.

- (b) 6 points Explain what the **competitive factor** is, and demonstrate that for EDF this equals to zero.

Solution: The competitive factor ϕ captures how well (bad) a scheduling policy A does in comparison to the optimal (offline) scheduler.

$$\phi_a = \min_T \frac{\Gamma_A(T)}{\Gamma_{max}(T)}$$

For EDF we can construct an example with just two tasks of value ϵ and 1, where the deadline of the former is earlier than the latter, and the unit task has no slack (deadline = computation time). See lecture notes.